



# AZ-203.3

## Module 02: Develop solutions that use Azure Cosmos DB

Prashanth Kumar



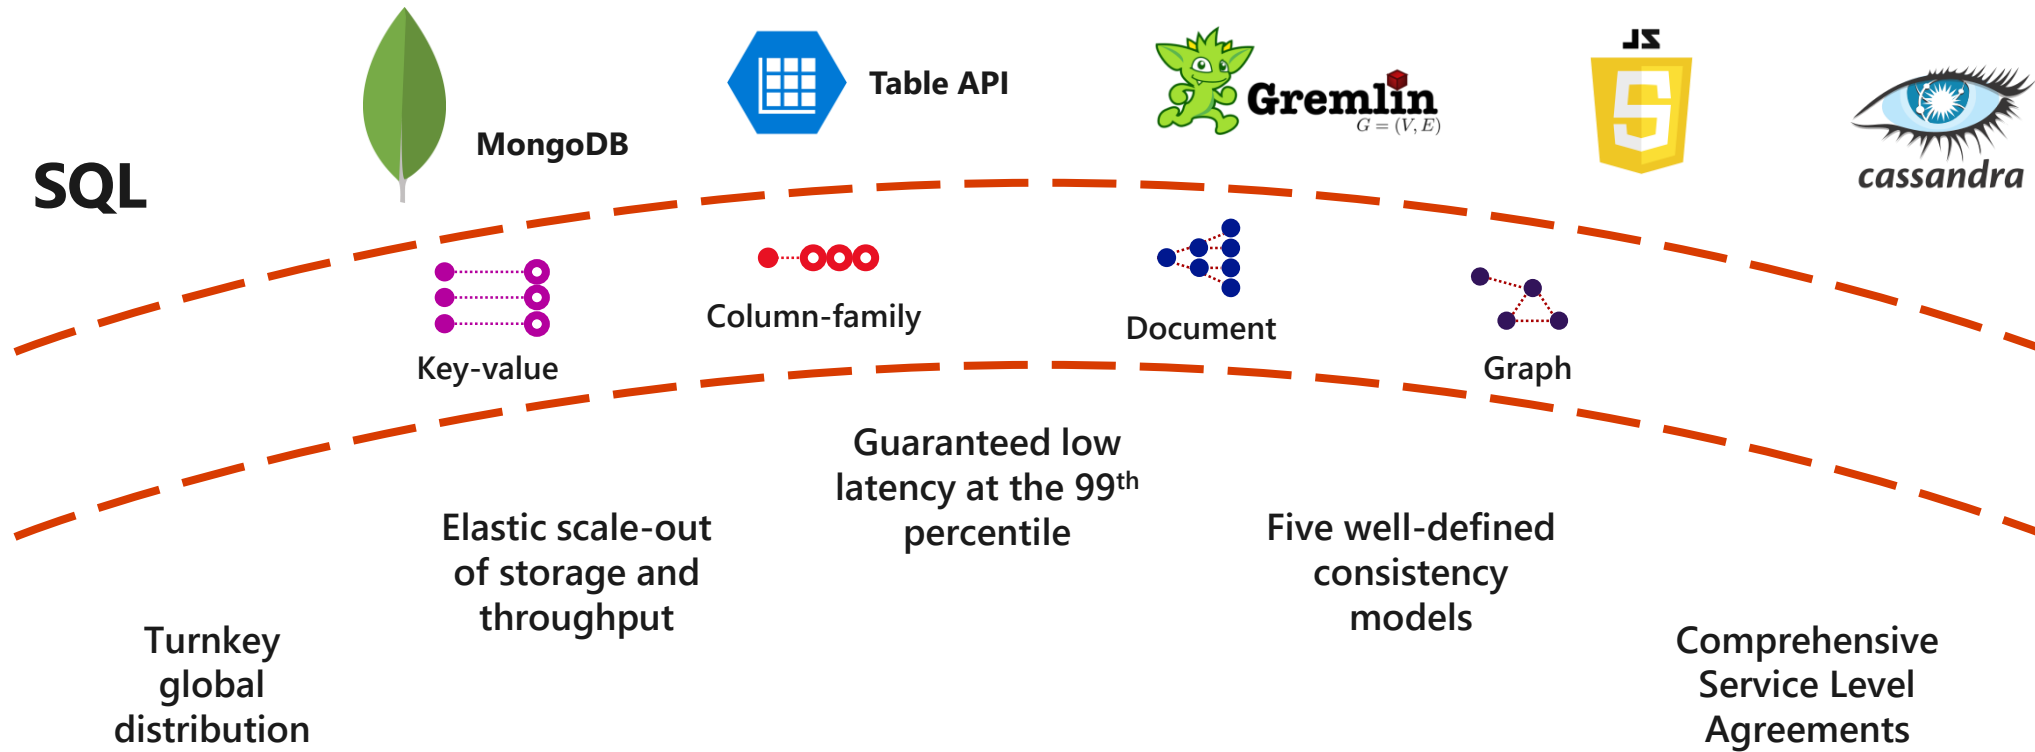
# Topics

- Azure Cosmos DB
- Managing containers and items
- Create and update documents by using code
- Server-side programming and features

# Lesson 01: Azure Cosmos DB




# Azure Cosmos DB




# Core functionality

- Global replication
  - Automatic and synchronous multi-region replication
  - Supports automatic and manual failover
- Varied consistency levels
  - Offers five consistency models
  - Provides control over performance-consistency tradeoffs, backed by comprehensive SLAs
- Low latency
  - Serve <10 ms read and <15 ms write requests at the 99th percentile
- Elastic scale-out
  - Elastically scale throughput from 10 to 100s of millions of requests/sec across multiple regions
  - Support for requests/sec for different workloads

# Global Replication




*Turnkey global distribution* automatically replicates data to other Azure datacenters across the globe without the need to manually write code or build a replication infrastructure



# Consistency levels

Azure Cosmos DB provides five consistency levels:



# Consistency levels

Consistency Level	Description
<b>Strong</b>	When a write operation is performed on your primary database, the write operation is replicated to the replica instances. The write operation is committed (and visible) on the primary only after it has been committed and confirmed by all replicas.
<b>Bounded Stateless</b>	This level is similar to the Strong level with the major difference that you can configure how stale documents can be within replicas. Staleness refers to the quantity of time (or the version count) a replica document can be behind the primary document.
<b>Session</b>	This level guarantees that all read and write operations are consistent within a user session. Within the user session, all reads and writes are monotonic and guaranteed to be consistent across primary and replica instances.
<b>Consistent Prefix</b>	This level has loose consistency but guarantees that when updates show up in replicas, they will show up in the correct order (that is, as prefixes of other updates) without any gaps.
<b>Eventual</b>	This level has the loosest consistency and essentially commits any write operation against the primary immediately. Replica transactions are asynchronously handled and will eventually (over time) be consistent with the primary. This tier has the best performance, because the primary database does not need to wait for replicas to commit to finalize its transactions.



# APIs



- MongoDB API

- Acts as a massively scalable MongoDB service powered by the Azure Cosmos DB platform
- Compatible with existing MongoDB libraries, drivers, tools, and applications



- Table API


- A key-value database service built to provide premium capabilities to existing Azure Table storage applications without making any app changes



- Gremlin API

- A fully managed, horizontally scalable graph database service
- Easy-to-build and run applications that work with highly connected datasets supporting Open Graph APIs (based on the Apache TinkerPop specification, Apache Gremlin)

# APIs (cont.)



- Cassandra API

- Globally distributed Apache Cassandra service powered by the Azure Cosmos DB platform
- Compatible with existing Apache Cassandra libraries, drivers, tools, and applications

- SQL API



- JavaScript and JavaScript Object Notation (JSON) native API based on the Azure Cosmos DB database engine
- Provides query capabilities rooted in SQL
- Query for documents based on their identifiers or make deeper queries based on properties of the document, complex objects, or the existence of specific properties
- Supports the execution of JavaScript logic within the database in the form of stored procedures, triggers, and user-defined functions

# Demo: Creating an Azure Cosmos DB account



# Consistency levels and Azure Cosmos DB APIs

## Apache Cassandra

Apache Cassandra 4.x	Azure Cosmos DB (multi-region)	Azure Cosmos DB (single region)
<b>ONE, TWO, THREE</b>	Consistent prefix	Consistent prefix
<b>LOCAL_ONE</b>	Consistent prefix	Consistent prefix
<b>QUORUM, ALL, SERIAL</b>	Bounded stateless	Strong
<b>LOCAL_QUORUM</b>	Bounded stateless	Strong
<b>LOCAL_SERIAL</b>	Bounded stateless	Strong

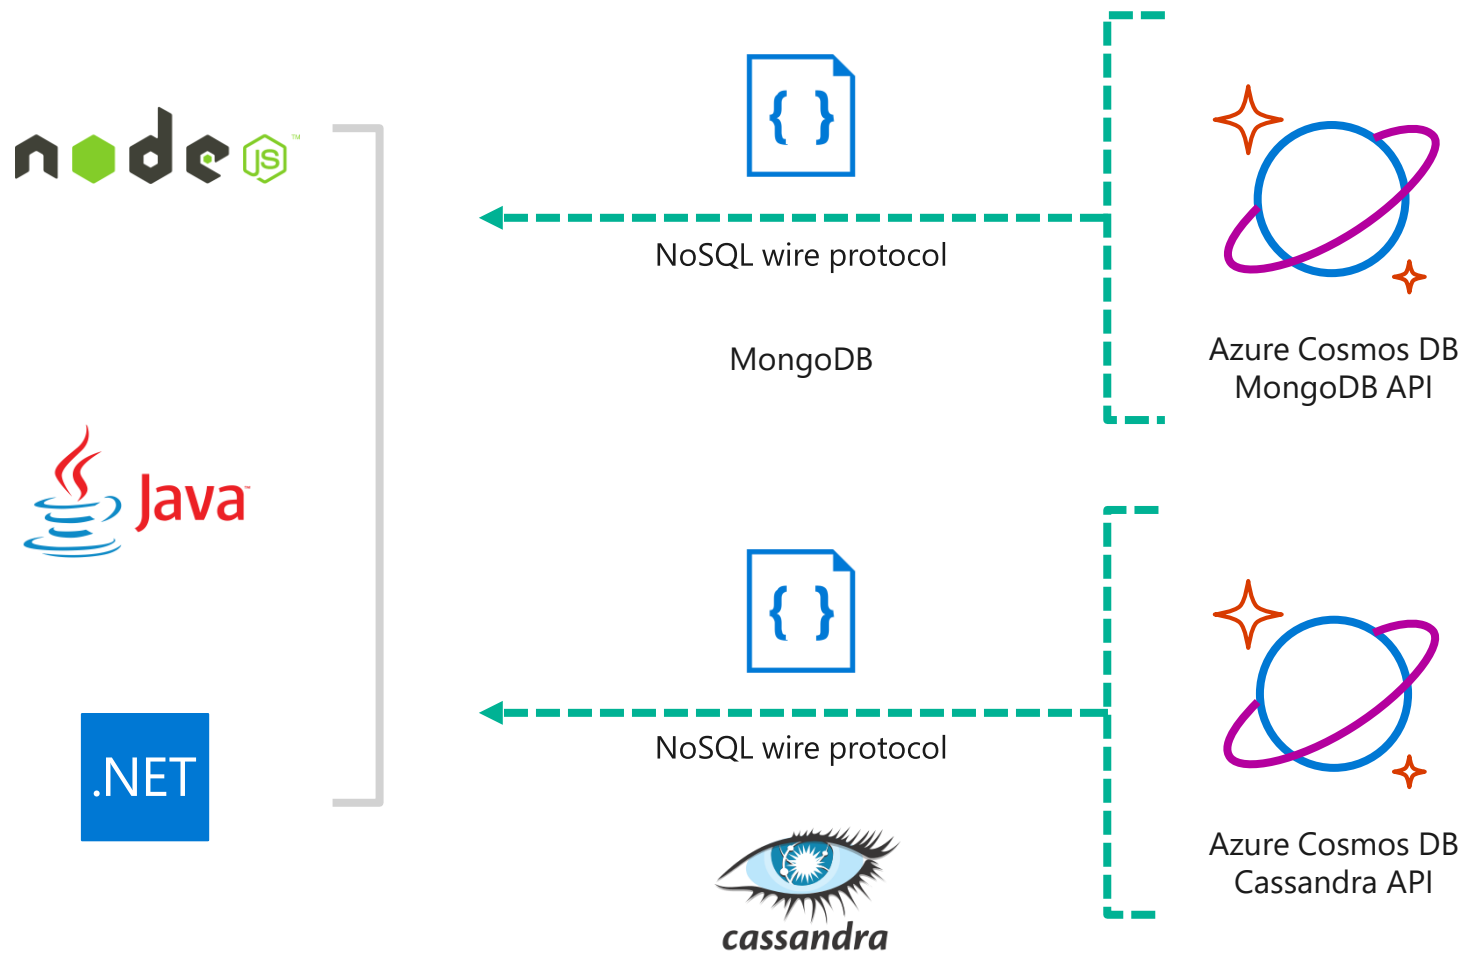
## MongoDB (3.4)

MongoDB 3.4	Azure Cosmos DB (multi-region)	Azure Cosmos DB (single region)
<b>Linearizable</b>	Strong	Strong
<b>Majority</b>	Bounded staleness	Strong
<b>Local</b>	Consistent prefix	Consistent prefix

# Migrating from NoSQL

- Many NoSQL database engines are simple to get started with, but they might cause problems as you scale, including:
  - Tedious setup and maintenance requirements for a multiple-server database cluster
  - Expensive and complex high-availability solutions
  - Challenges in achieving end-to-end security, including encryption at rest and in flight
  - Required resource overprovisioning and unpredictable costs to achieve scale
- Azure Cosmos DB provides NoSQL-as-a-service for:
  - MongoDB
  - Cassandra
  - Gremlin


# Migrating from NoSQL



# Lesson 02: Managing containers and items



# Resource hierarchy






# Resource hierarchy (continued)

Resource	Description
<b>Account</b>	A set of databases
<b>Database</b>	Logical container for containers that can (optionally) share throughput across the containers
<b>Collection (container)</b>	A group of Items and programmatic resources usually related in some way
<b>Document (item)</b>	An arbitrary unit of content In many cases, this would be a JSON document
<b>Stored procedure (sproc)</b>	Application logic written in JavaScript executed within the database engine as a transaction
<b>Trigger</b>	Application logic written in JavaScript executed before or after either an insert, replace, or delete operation
<b>User-defined function</b>	Application logic written in JavaScript to extend the SQL API query language

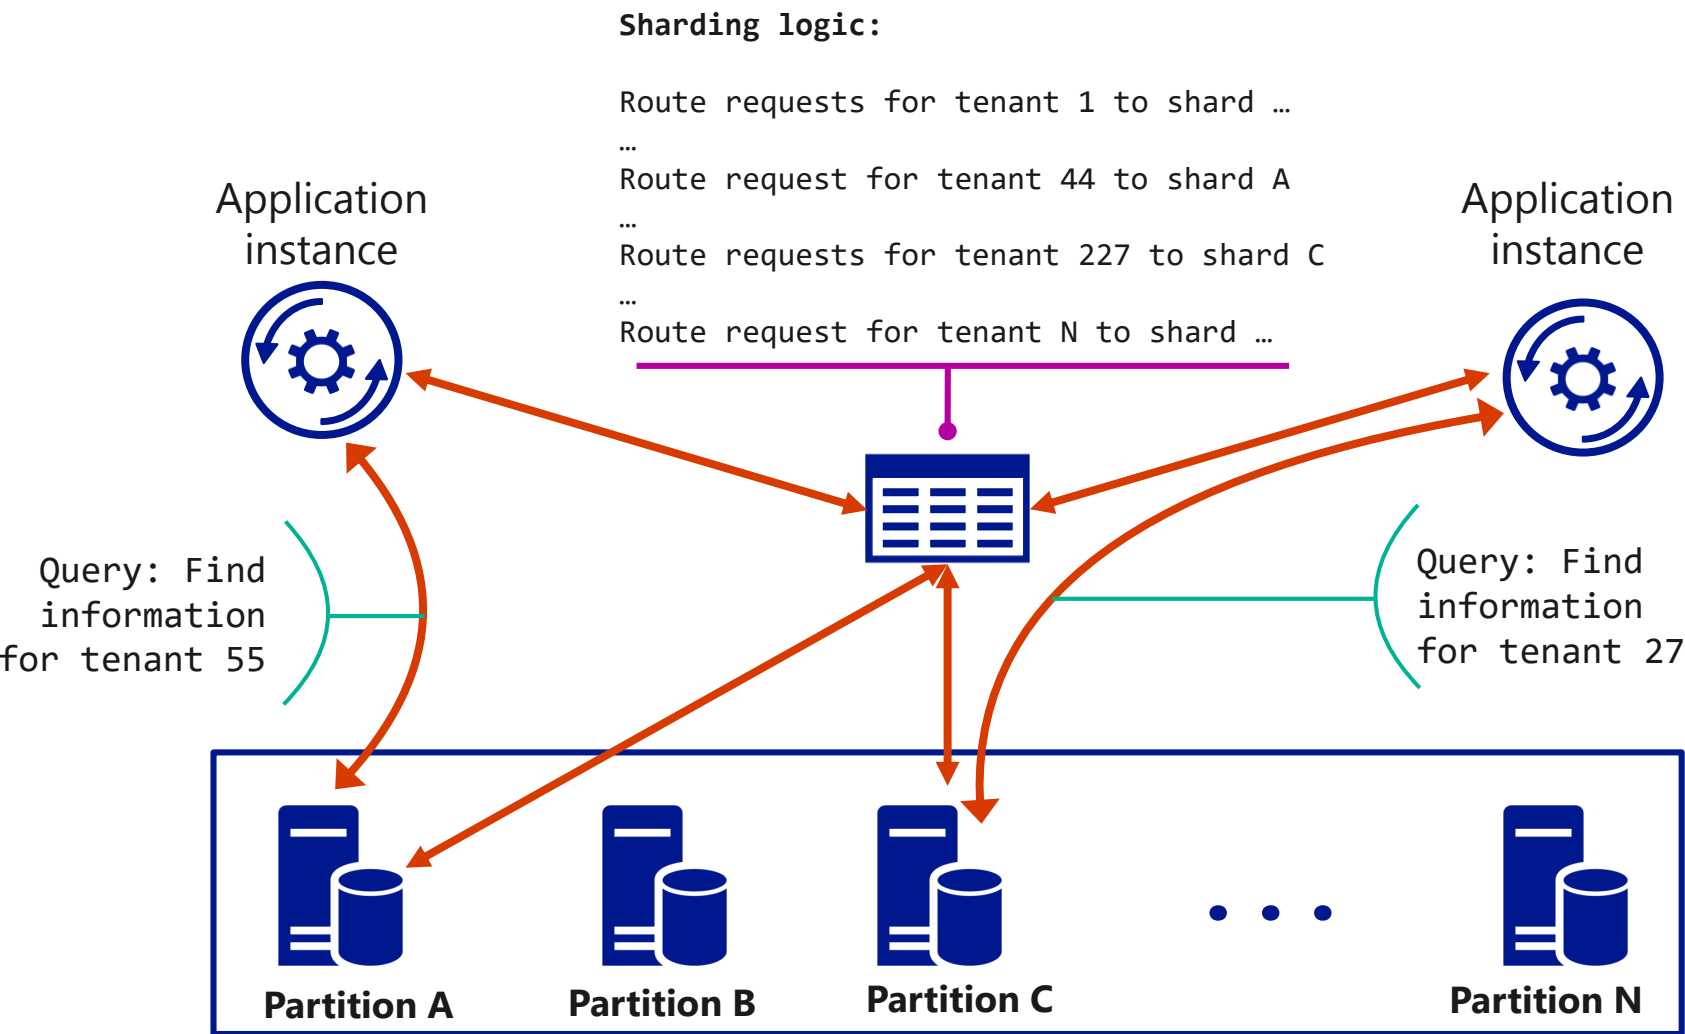
# Containers



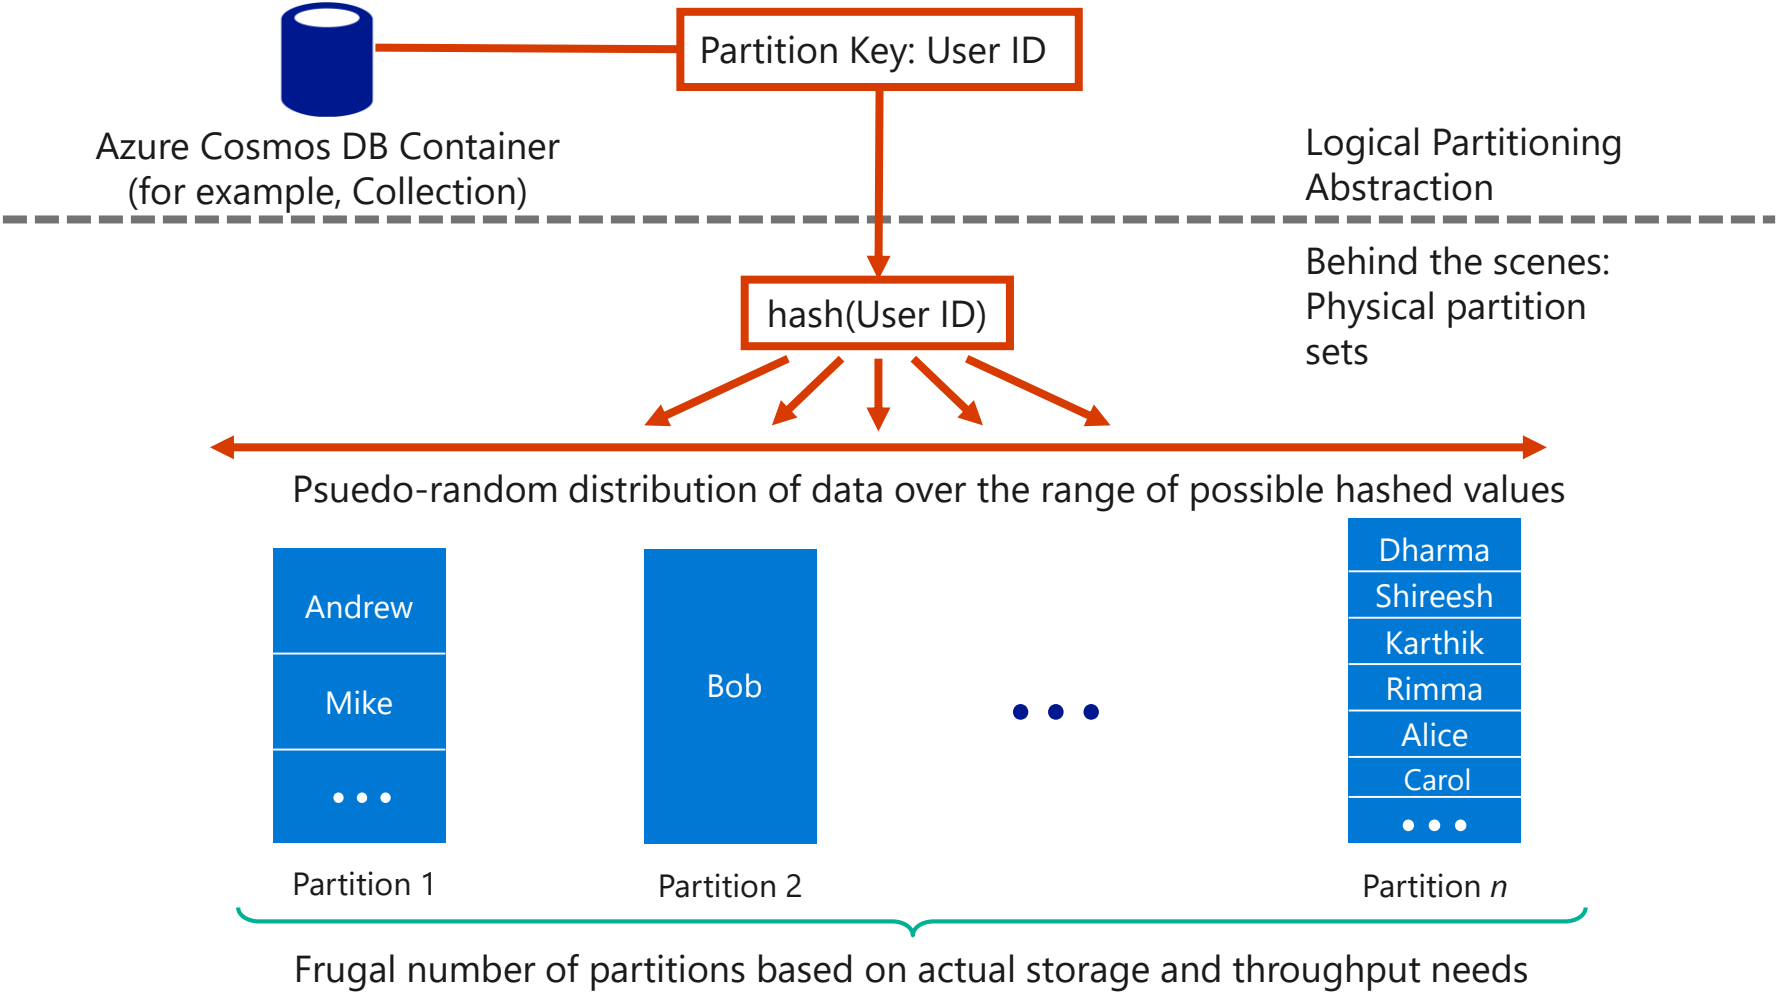
# Demo: Dynamically adjusting container throughput



# Partitioning



# Partitioning implementation




# Lesson 03: Create and update documents by using code



# Manage collections and documents

- Install the **Microsoft.Azure.DocumentDB.Core** package from NuGet
- Use the following namespaces
  - Microsoft.Azure.Documents
  - Microsoft.Azure.Documents.Client
- Use the DocumentClient class



# Accessing collections by using .NET

```
using Microsoft.Azure.Documents;  
using Microsoft.Azure.Documents.Client;  
  
DocumentClient client = new DocumentClient(new Uri("[endpoint]"), "[key]");  
  
Uri collectionUri = UriFactory.CreateDocumentCollectionUri(databaseName,  
collectionName);
```





# Reading documents by using .NET

```
// create Uri referencing document
```

```
Uri documentUri = UriFactory.CreateDocumentUri(databaseName, collectionName,  
documentId);
```

```
// Use the ReadDocumentAsync method
```

```
SerializedType document = await this.client.  
ReadDocumentAsync<SerializedType>(documentUri);
```



# Creating documents by using .NET

```
var document = new {  
    firstName = "Alex",  
    lastName = "Leh"  
}  
  
await this.client.CreateDocumentAsync(collectionUri, document);
```



# Querying documents by using .NET (continued 1)

```
var query = client.CreateDocumentQuery<Family>(
    collectionUri, new SqlQuerySpec() {
        QueryText = "SELECT * FROM f WHERE (f.surname = @lastName)",
        Parameters = new SqlParameterCollection()
        {
            new SqlParameter("@lastName", "Andt")
        }
    }, DefaultOptions
);

var families = query.ToList();
```



# Querying documents by using .NET (continued 2)

```
var query = client.CreateDocumentQuery<Family>(collectionUri)
    .Where(d => d.Surname = "Andt")
    .Select(d => new { Name = d.Id, City = d.Address?.City})
    .AsDocumentQuery();

var families = query.ToList();
```



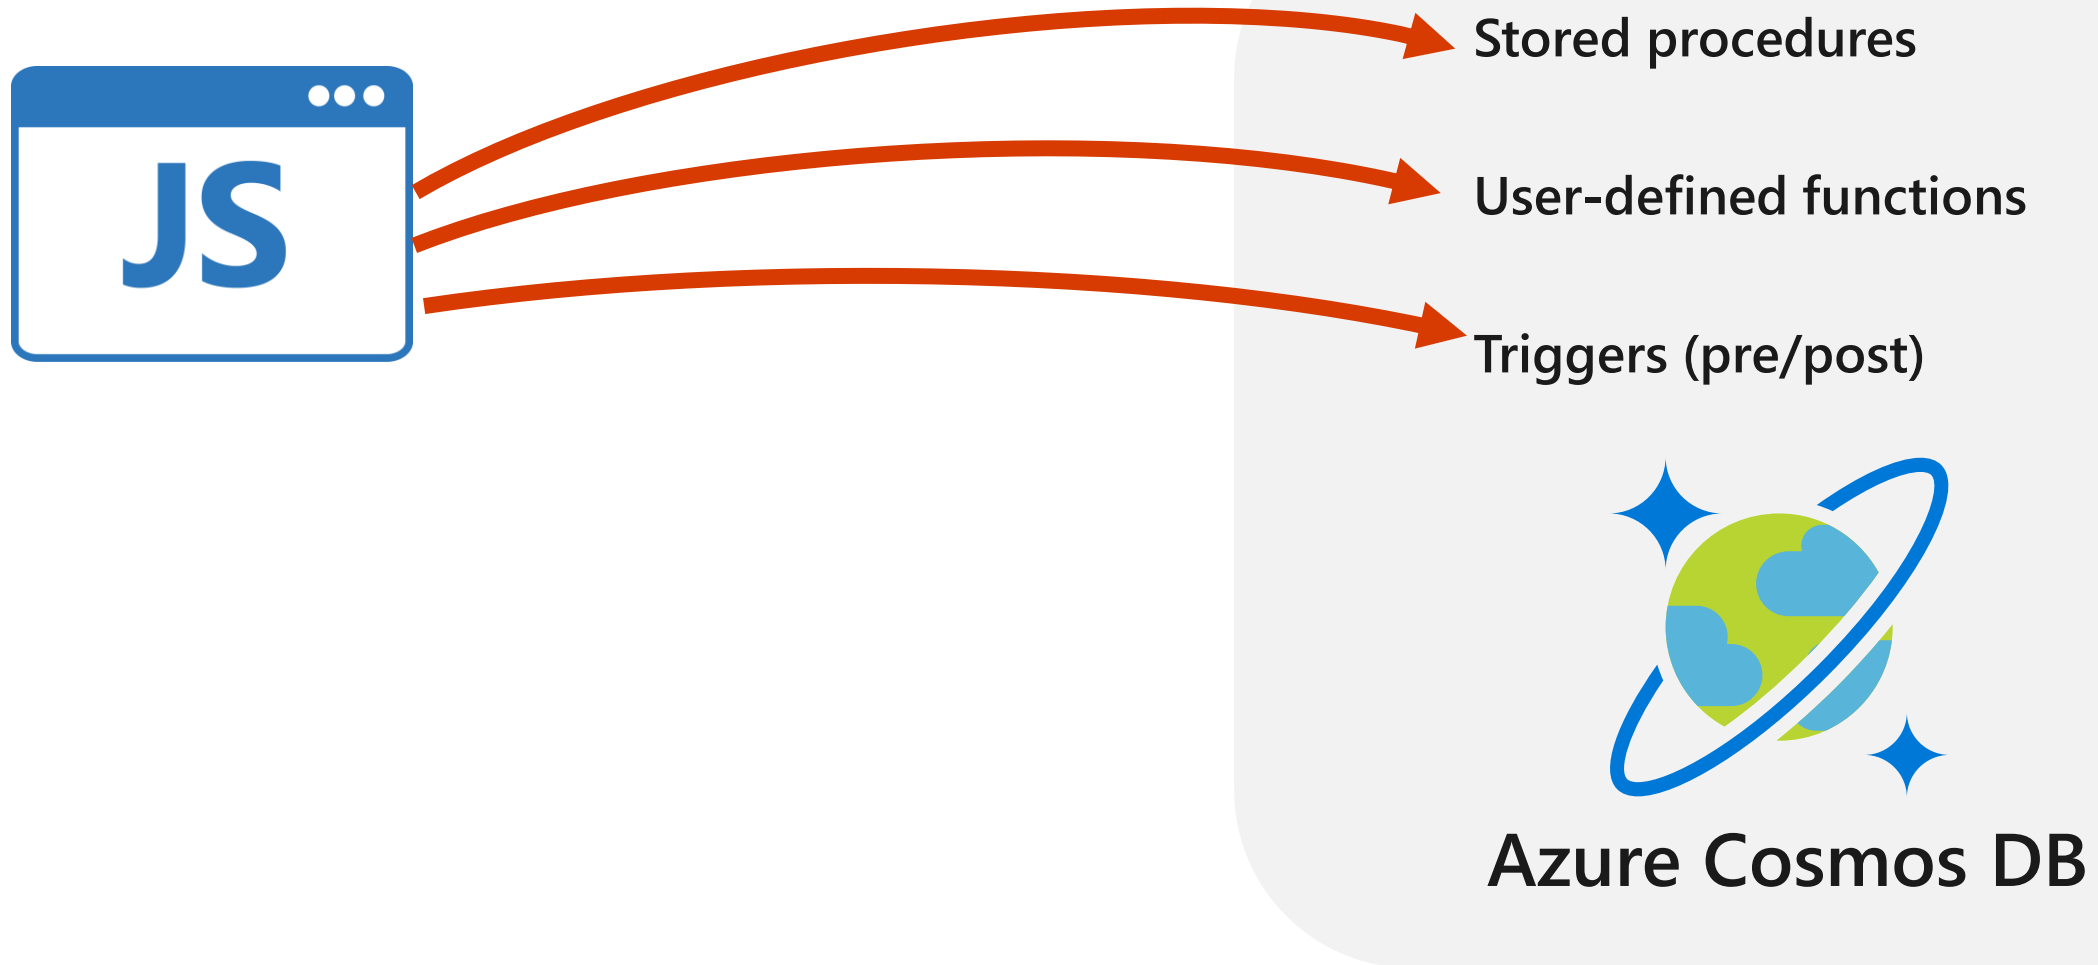
# Demo: Managing Azure Cosmos DB by using .NET



# Lesson 04: Server-side programming and features

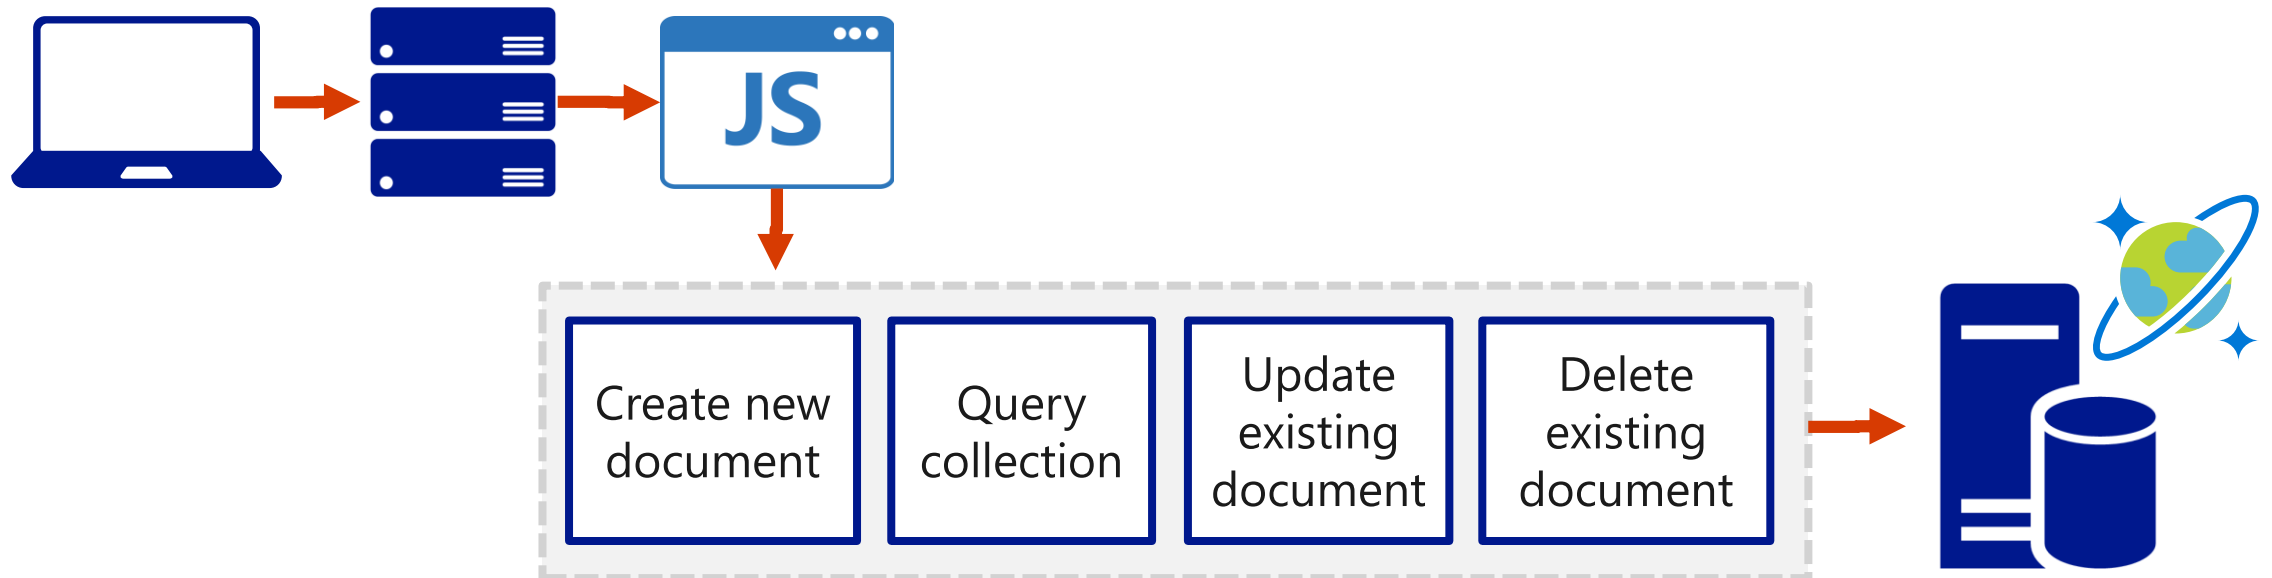


# JavaScript and Azure Cosmos DB



# Stored procedures

- In Azure Cosmos DB, JavaScript is hosted in the same memory space as the database
- Requests made within stored procedures and triggers run in the same scope of a database session






# Stored procedure in JavaScript

```
function createSampleDocument(documentToCreate) {  
    var context = getContext();  
    var collection = context.getCollection();  
    var accepted = collection.createDocument(  
        collection.getSelfLink(),  
        documentToCreate,  
        function (error, documentCreated) {  
            context.getResponse().setBody(documentCreated.id)  
        }  
    );  
    if (!accepted) return;  
}
```



# Bounded execution

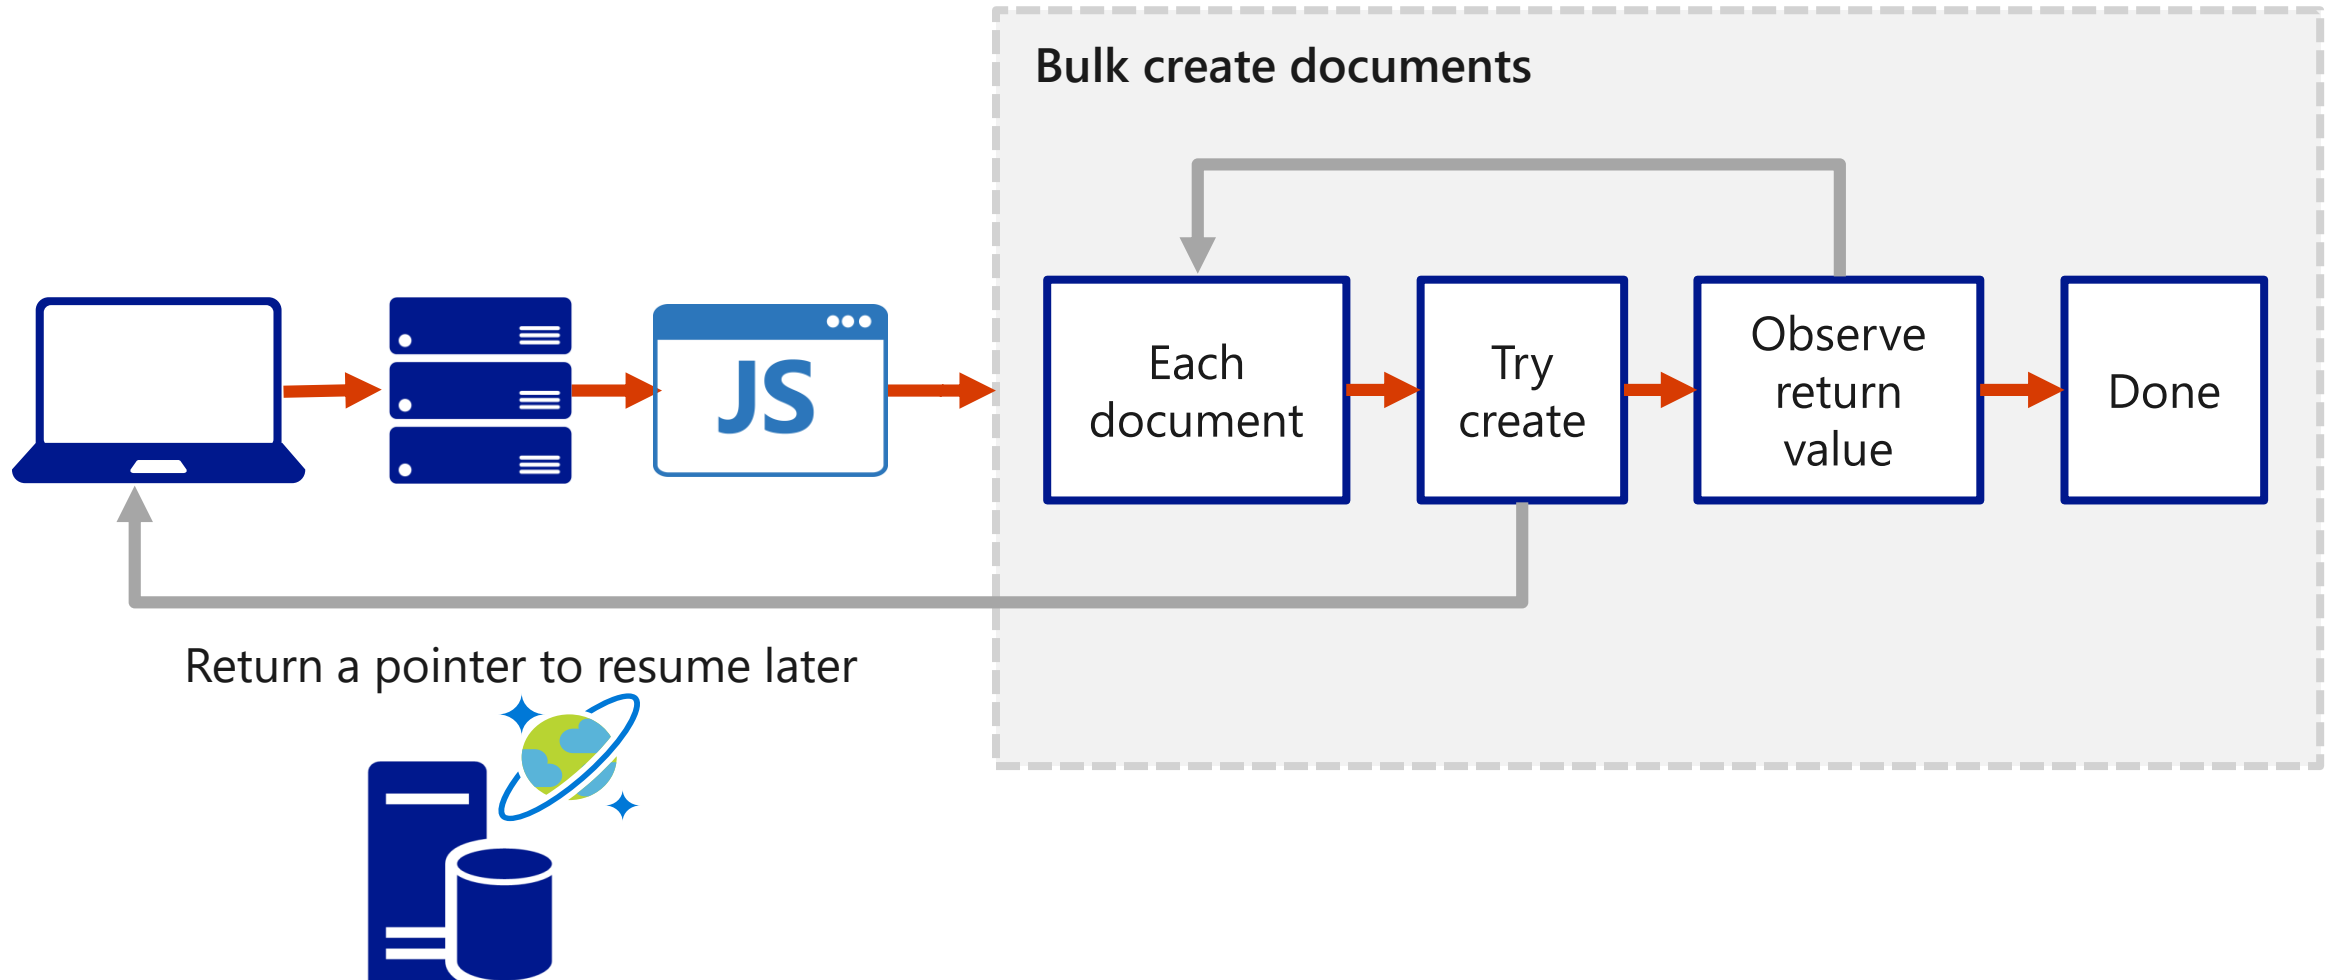
- All Azure Cosmos DB operations must complete within a limited amount of time
  - Specifically, stored procedures have a limited amount of time to run on the server
- All collection functions return a Boolean value that represents whether that operation will complete or not



# Transaction continuation

- JavaScript functions can implement a continuation-based model to batch or resume execution
- The continuation value can be any value of your choice
- Your applications can then use this value to resume a transaction from a new starting point

# Transaction continuation (cont.)



# User-defined functions in JavaScript

```
var taxUdf = {  
  id: "tax",  
  serverScript: function tax(income) {  
    if (income == undefined)  
      throw 'no input';  
    if (income < 1000)  
      return income * 0.1;  
    else if (income < 10000)  
      return income * 0.2;  
    else  
      return income * 0.4;  
  }  
}
```



# User-defined functions in SQL queries

SELECT

\*

FROM

TaxPayers t

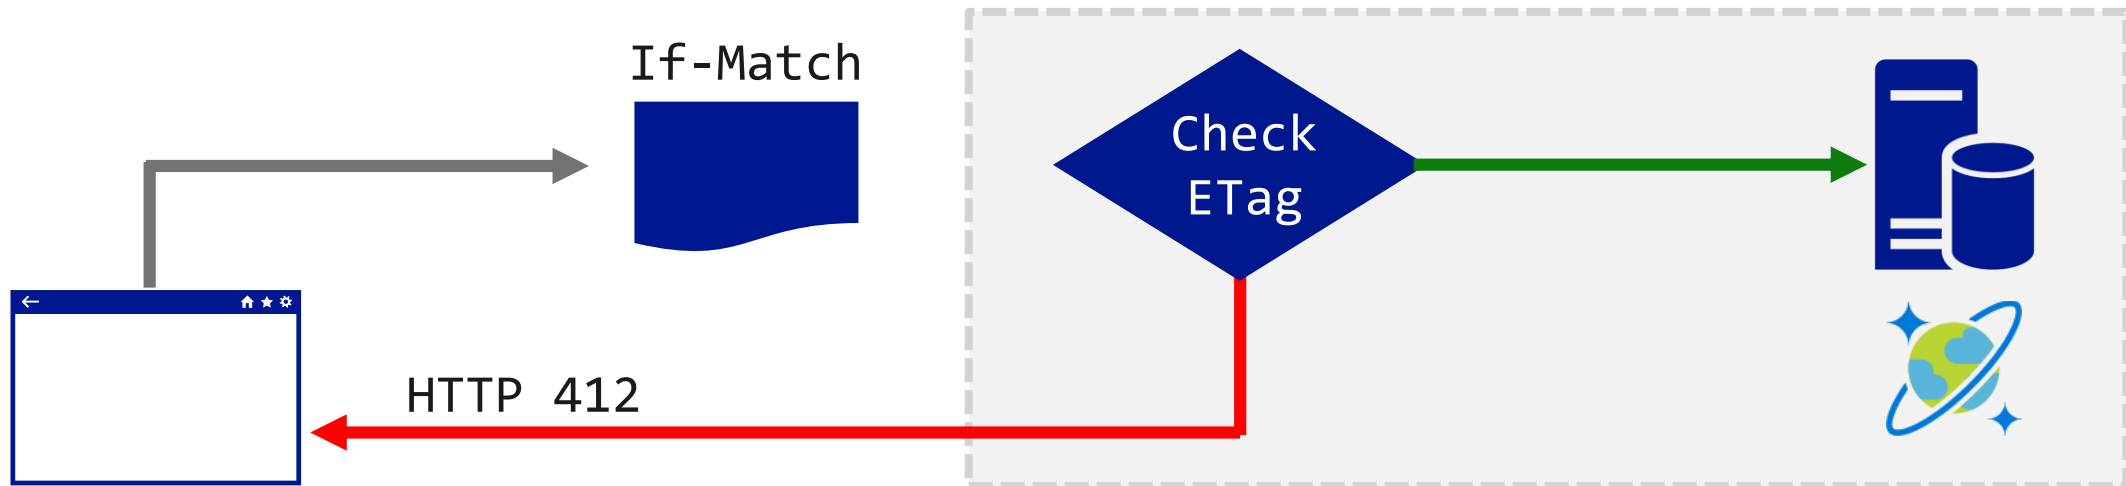
WHERE

udf.tax(t.income) > 20000



# Optimistic concurrency

- The SQL API supports optimistic concurrency control through HTTP ETags
- Every SQL API resource has an ETag system property
- ETags can be used with the If-Match HTTP request header to allow the server to decide whether a resource should be updated



# Controlling concurrency in .NET

```
try
{
    var ac = new AccessCondition { Condition = readDoc.ETag, Type =
        AccessConditionType.IfMatch };
    await client.ReplaceDocumentAsync(readDoc, new RequestOptions {
        AccessCondition = ac });
}
catch (DocumentClientException dce)
{
    if (dce.StatusCode == HttpStatusCode.PreconditionFailed)
    {
        Console.WriteLine("Another process has updated the record");
    }
}
```





# Review

- Azure Cosmos DB
- Managing containers and items
- Create and update documents by using code
- Server-side programming and features

